

Advanced JPEG Steganography and Detection

Overview:

- The JPEG Algorithm
- Exploiting the JPEG algorithm
- Detecting Exploited JPEGs

Why JPEG? - The JPEG Algorithm

- The Joint Photographic Experts Group (JPEG) algorithm combines both lossless and lossy data compression techniques
- Achieves a small file size with little perceptible difference
- Particularly suited to natural images
- One of the most common image formats if not the most common

Algorithm Overview

Colour Plane Conversion → Convert unsigned to signed → DCT → Quantiser → Run-length Encoding → Entropy Encoder

Colour Plane Conversion

- The JPEG algorithm first converts RGB to YC_rC_b
 - Y is the luminance component
 - C_r and C_b are the red difference and blue difference chroma components
 - Grayscale images only have the Y component
- The human eye is less sensitive to chrominance than luminance

Discrete Cosine Transform

- The 2-dimensional Discrete Transform is applied to an 8×8 image block
- This process breaks down the frequency components of an image

- Low frequencies are smooth transitions in an image
- High Frequencies are sudden changes (like in a cartoon)
- The purpose of this is to modulate the influence of different spectral components on the image
 - i.e. Higher Frequencies contribute less information to the image and therefore can be reduced or eliminated

Quantisation and Entropy Coding

- The DCT results are quantised at a desired quality level
- Entropy coding is applied
 - A combination of Run-Length Encoding and Huffman coding
- To view the image, the process is reversed
- The restored image looks similar or almost exact to the human, but mathematically it is completely different
- For a high quality JPEG, there is almost no perceptible difference

Quantisation and Coding

- An 8×8 quantisation table is used to scale these coefficients
 - This is where the greatest loss occurs
 - The result is the "quantised DCT coefficients"
 - The Q-Table controls the quality
- The quantised DCT coefficient matrix generally has a lot of ZERO values, which are Run-Length coded away
 - The remaining compression is lossless
- With the small remaining numbers, Huffman or Arithmetic encoding is applied
- Loop to the next 8×8 block and repeat until the image is complete

Exploiting the JPEG Algorithm - JPEG Hiding

Swap DCT

- Choose two DCT coefficients that have the same value in the quantisation table
 - Select middle frequencies so hidden bits are in significant portions of the image
- Select a cover block
- Get DCT transform of the block
- Read a message bit from the file to be hidden
- If the condition is already true, continue
- If the condition is not true, SWAP the coefficients
 - Note: this is done prior to quantisation, so the difference must be large enough to hold true after quantisation

Weaknesses in this approach

- A particular cover block may be a poor candidate for hiding
- Capacity is 1 bit per 8×8 block
 - For a 256×256 image, that's 32×32 = 1024 blocks max

DCT Least Significant Bit (LSB)

- JSteg uses this approach
- Alter the LSB of each quantised DCT coefficient to hold our message
- More than one bit/block capacity
 - Depends on the number of non-zero coefficients
- Can use any coefficients that are not zero or one
 - Using zero would increase capacity, but also distortion, and image size
 - A clear indication of data hiding would be a low number of DCT coefficients with a zero value
- Can't use '1' because ... 0001₂ → change LSB and 0000₂

Outguess

- Hides in LSBs of DCT coefficients

- Pseudo-randomly permutes selection DCT coefficients that are not zero or one
- After embedding, a second pass is made to make sure corrections to unused coefficients, such that DCT histogram is preserved
 - Reduces capacity as some coefficients are used for correction
 - Makes detection more difficult

F5

- F5 takes a different approach to hiding in the DCT coefficients
- F5 has a fairly high capacity, but very low detectability
- F5 decrements the magnitude of the coefficient values when the LSB does not match the message
 - As opposed to overwhelming them with the message bits
 - Note that 1 and -1 become zero - called shrinkage
 - must be decremented to zero since a 2 will become a 1
- Skips a zero for embedding and extraction
- Inverts the meaning for negative DCT coefficients
 - An LSB of 1 in a negative coefficient represents a zero
 - Prevents uneven distribution of odd vs even coefficients
- Uses permutative straddling
 - Spreads the message over the entire image
 - Like the cryptographic spreading of the other techniques

Statistically Invisible Steganography

- SIS performs a complexity analysis of each 8×8 DCT block
- Number of non-zero coefficients must exceed a threshold or the entire block is skipped
 - threshold = 0.3 to 0.6
 - 20 to 29 coefficients out of a block must be non-zero
- Adds up different sets of coefficients to produce a sum

- If the LSB of the sum equals the messages, next block
- If not, add/subtract 1 from the largest magnitude

YASS

- Yet another Steganographic Scheme that resists blind steganalysis
- What YASS does a little differently is to select blocks larger than 8×8
 - Example: 10×10 blocks
 - Has 9 possible sub-blocks
- Out of the larger block, YASS selects an 8×8 block, performs the DCT conversion and quantisation
- Hides in those coefficients
- Must use an error correcting code since there will be some errors when converted to JPEG

High-Capacity DCT

- An adaptive DCT LSB technique
 - Hides mostly in lower and middle frequency components
 - Performs a capacity estimation
 - Adapts to different characteristics of each block
- Experimentally, the quality must be $\geq \sim 75$ to remain imperceptible
 - At quality = 50, (the standard) visual distortion is obvious
 - At quality = 60, it is noticeable if you are looking for it

Steganalysis - Detecting Exploited JPEGs

- Three levels of defeat for steganography
 - Detection
 - Extraction
 - Destruction
- Detection ability based on embedded data size

- If you embed a single bit into a single DCT coefficient, it is much less detectable

General Approach

- Get as much information as possible
 - Adversary's goal
 - Tool(s) likely used
 - Types of cover files
 - Type of message(s)
- Check for tool signatures
 - Ex: modified Quantisation table
 - Specific file existence
 - If you are forensically examining a disk
 - Specific types of distortion
 - For JPEG, the typical artifact is blockiness
- Type of cover files
- Type of message(s)
 - Text
 - Encrypted/compressed
 - Other images
- Apply analysis specific to the tool or to the target cover files

Defeating JPEG Steganography

- Two general approaches to JPEG Detection
 - Analysis of DCT coefficients
 - Block edge detection
 - Blockiness
- Some techniques rely on training on clean images
- Apply Statistics

- Histograms and Entropy
- Chi-Squared Test

Weaknesses:

- Outguess uses excess capacity to make adjustments to DCT coefficients
 - Keeps the balance
- SwapDCT does not change the value of any coefficients
 - This analysis reveals nothing for swapDCT
- F5 mitigates changes in coefficients
 - Uses matrix encoding to reduce the actual number of changes
 - Does not substitute bits, decrements existing values, maintaining balance
- For these and other techniques, different detection methods are needed

Workarounds:

- An approach to detect F5 is to predict the histogram of the original cover image
 - F5 does increase the number of zero coefficients
- Decompress the stego image, crop it by 4 columns, and recompress using the same quantisation table
 - Spatially, an image cropped by just 4 vertical columns is nearly identical
- Apply a blurring algorithm to reduce blockiness introduced by the cropping
- Compare a predicted histogram with the stego-image histogram
- Able to calculate the approximate message length as well

Extracting JPEG Steganography

- Extraction is much more difficult than detection
- Cryptography complicates extraction
 - Doesn't prevent detection
- Knowing the method is critical

- If you extract LSBs from a JPEG that used Swap DCT, you gain no information about the message

Destroying JPEG Steganography

- Sterilisation of data hidden in a jpeg is easy
- Could ZERO or RANDOMISE the LSBs of the DCT coefficients
 - Too hard in some instances
- Could hide another message on top of prior message
 - Similar to randomisation
 - Use the same tool if known
- Resize the iamge
 - Not in multiples of 8
 - Resize by a single horizontal column and vertical rows
 - Completely changes DCT coefficients